

bonus__merge

February 29, 2024

1 Bonus video: How do I merge DataFrames in pandas?

Full course: [pandas in 30 days](#)

© 2024 Data School. All rights reserved.

1. Selecting a Function
2. Joining (Merging) DataFrames
3. What if...?
4. Four Types of Joins

1.1 Part 1: Selecting a Function

Taken from “Merging DataFrames with pandas” (DataCamp course):

- `df1.append(df2)`: stacking vertically
 - Update from 2024: The “append” method has been removed in favor of the `concat` function.
- `pd.concat([df1, df2])`:
 - stacking many horizontally or vertically
 - simple inner/outer joins on Indexes
- `df1.join(df2)`: inner/outer/left/right joins on Indexes
- `pd.merge(df1, df2)`: many joins on multiple columns

1.2 Part 2: Joining (Merging) DataFrames

Using the [MovieLens 100k data](#), let’s create two DataFrames:

- **movies**: shows information about movies, namely a unique **movie_id** and its **title**
- **ratings**: shows the **rating** that a particular **user_id** gave to a particular **movie_id** at a particular **timestamp**

```
[1]: import pandas as pd
```

1.2.1 Movies

```
[2]: movie_cols = ['movie_id', 'title']
movies = pd.read_table('http://bit.ly/movieitems', sep='|', header=None,
    ↪ names=movie_cols, usecols=[0, 1])
movies.head()
```

```
[2]:      movie_id      title
      0         1  Toy Story (1995)
      1         2  GoldenEye (1995)
      2         3  Four Rooms (1995)
      3         4  Get Shorty (1995)
      4         5   Copycat (1995)
```

```
[3]: movies.shape
```

```
[3]: (1682, 2)
```

```
[4]: movies.movie_id.nunique()
```

```
[4]: 1682
```

1.2.2 Ratings

```
[5]: rating_cols = ['user_id', 'movie_id', 'rating', 'timestamp']
      ratings = pd.read_table('http://bit.ly/movielensdata', sep='\t', header=None,
      ↪names=rating_cols)
      ratings.head()
```

```
[5]:      user_id  movie_id  rating  timestamp
      0       196       242       3  881250949
      1       186       302       3  891717742
      2        22       377       1  878887116
      3       244        51       2  880606923
      4       166       346       1  886397596
```

```
[6]: ratings.shape
```

```
[6]: (100000, 4)
```

```
[7]: ratings.movie_id.nunique()
```

```
[7]: 1682
```

```
[8]: ratings.loc[ratings.movie_id == 1, :].head()
```

```
[8]:      user_id  movie_id  rating  timestamp
      24       308         1         4  887736532
      454       287         1         5  875334088
      957       148         1         4  877019411
      971       280         1         4  891700426
     1324        66         1         3  883601324
```

1.2.3 Merging Movies and Ratings

Let's pretend that you want to examine the ratings DataFrame, but you want to know the **title** of each movie rather than its **movie_id**. The best way to accomplish this objective is by “joining” (or “merging”) the DataFrames using the pandas `merge` function:

```
[9]: movies.columns
[9]: Index(['movie_id', 'title'], dtype='object')

[10]: ratings.columns
[10]: Index(['user_id', 'movie_id', 'rating', 'timestamp'], dtype='object')

[11]: movie_ratings = pd.merge(movies, ratings)
      movie_ratings.columns
[11]: Index(['movie_id', 'title', 'user_id', 'rating', 'timestamp'], dtype='object')

[12]: movie_ratings.head()
[12]:
```

	movie_id	title	user_id	rating	timestamp
0	1	Toy Story (1995)	308	4	887736532
1	1	Toy Story (1995)	287	5	875334088
2	1	Toy Story (1995)	148	4	877019411
3	1	Toy Story (1995)	280	4	891700426
4	1	Toy Story (1995)	66	3	883601324

```
[13]: movie_ratings.shape
[13]: (100000, 5)
```

Here's what just happened:

- pandas noticed that movies and ratings had one column in common, namely **movie_id**. This is the “key” on which the DataFrames will be joined.
- The first **movie_id** in movies is 1. Thus, pandas looked through every row in the ratings DataFrame, searching for a **movie_id** of 1. Every time it found such a row, it recorded the **user_id**, **rating**, and **timestamp** listed in that row. In this case, it found 452 matching rows.
- The second **movie_id** in movies is 2. Again, pandas did a search of ratings and found 131 matching rows.
- This process was repeated for all of the remaining rows in movies.

At the end of the process, the `movie_ratings` DataFrame is created, which contains the two columns from movies (**movie_id** and **title**) and the three other columns from ratings (**user_id**, **rating**, and **timestamp**).

- **movie_id** 1 and its **title** are listed 452 times, next to the **user_id**, **rating**, and **timestamp** for each of the 452 matching ratings.

- **movie_id** 2 and its **title** are listed 131 times, next to the **user_id**, **rating**, and **timestamp** for each of the 131 matching ratings.
- And so on, for every movie in the dataset.

```
[14]: print(movies.shape)
      print(ratings.shape)
      print(movie_ratings.shape)
```

```
(1682, 2)
(100000, 4)
(100000, 5)
```

Notice the shapes of the three DataFrames:

- There are 1682 rows in the movies DataFrame.
- There are 100000 rows in the ratings DataFrame.
- The `merge` function resulted in a `movie_ratings` DataFrame with 100000 rows, because every row from ratings matched a row from movies.
- The `movie_ratings` DataFrame has 5 columns, namely the 2 columns from movies, plus the 4 columns from ratings, minus the 1 column in common.

By default, the `merge` function joins the DataFrames using all column names that are in common (`movie_id`, in this case). The [documentation](#) explains how you can override this behavior.

1.3 Part 3: What if...?

1.3.1 What if the columns you want to join on don't have the same name?

```
[15]: movies.columns = ['m_id', 'title']
      movies.columns
```

```
[15]: Index(['m_id', 'title'], dtype='object')
```

```
[16]: ratings.columns
```

```
[16]: Index(['user_id', 'movie_id', 'rating', 'timestamp'], dtype='object')
```

```
[17]: pd.merge(movies, ratings, left_on='m_id', right_on='movie_id').head()
```

```
[17]:   m_id      title  user_id  movie_id  rating  timestamp
0     1  Toy Story (1995)    308         1         4  887736532
1     1  Toy Story (1995)    287         1         5  875334088
2     1  Toy Story (1995)    148         1         4  877019411
3     1  Toy Story (1995)    280         1         4  891700426
4     1  Toy Story (1995)     66         1         3  883601324
```

1.3.2 What if you want to join on one index?

```
[18]: movies = movies.set_index('m_id')
      movies.head()
```

```
[18]:
```

	m_id	title
1	1	Toy Story (1995)
2	2	GoldenEye (1995)
3	3	Four Rooms (1995)
4	4	Get Shorty (1995)
5	5	Copycat (1995)

```
[19]: pd.merge(movies, ratings, left_index=True, right_on='movie_id').head()
```

```
[19]:
```

		title	user_id	movie_id	rating	timestamp
24	1	Toy Story (1995)	308	1	4	887736532
454	1	Toy Story (1995)	287	1	5	875334088
957	1	Toy Story (1995)	148	1	4	877019411
971	1	Toy Story (1995)	280	1	4	891700426
1324	1	Toy Story (1995)	66	1	3	883601324

1.3.3 What if you want to join on two indexes?

```
[20]: ratings = ratings.set_index('movie_id')
      ratings.head()
```

```
[20]:
```

	movie_id	user_id	rating	timestamp
242	196	3	881250949	
302	186	3	891717742	
377	22	1	878887116	
51	244	2	880606923	
346	166	1	886397596	

```
[21]: pd.merge(movies, ratings, left_index=True, right_index=True).head()
```

```
[21]:
```

	m_id	title	user_id	rating	timestamp
1	1	Toy Story (1995)	308	4	887736532
1	1	Toy Story (1995)	287	5	875334088
1	1	Toy Story (1995)	148	4	877019411
1	1	Toy Story (1995)	280	4	891700426
1	1	Toy Story (1995)	66	3	883601324

1.4 Part 4: Four Types of Joins

There are actually four types of joins supported by the pandas `merge` function. Here's how they are described by the [documentation](#):

- **inner:** use intersection of keys from both frames, similar to a SQL inner join; preserve the order of the left keys
- **outer:** use union of keys from both frames, similar to a SQL full outer join; sort keys lexicographically
- **left:** use only keys from left frame, similar to a SQL left outer join; preserve key order
- **right:** use only keys from right frame, similar to a SQL right outer join; preserve key order

The default is the “inner join”, which was used when creating the `movie_ratings` DataFrame.

It’s easiest to understand the different types by looking at some simple examples:

1.4.1 Example DataFrames A and B

```
[22]: A = pd.DataFrame({'color': ['green', 'yellow', 'red'], 'num':[1, 2, 3]})
A
```

```
[22]:      color  num
0   green    1
1  yellow    2
2    red     3
```

```
[23]: B = pd.DataFrame({'color': ['green', 'yellow', 'pink'], 'size':['S', 'M', 'L']})
B
```

```
[23]:      color size
0   green    S
1  yellow    M
2   pink    L
```

1.4.2 Inner join

Only include observations found in both A and B:

```
[24]: pd.merge(A, B, how='inner')
```

```
[24]:      color  num size
0   green    1    S
1  yellow    2    M
```

1.4.3 Outer join

Include observations found in either A or B:

```
[25]: pd.merge(A, B, how='outer')
```

```
[25]:      color  num size
0   green  1.0    S
1  yellow  2.0    M
2    red   3.0  NaN
3   pink  NaN    L
```

1.4.4 Left join

Include all observations found in A:

```
[26]: pd.merge(A, B, how='left')
```

```
[26]:
```

	color	num	size
0	green	1	S
1	yellow	2	M
2	red	3	NaN

1.4.5 Right join

Include all observations found in B:

```
[27]: pd.merge(A, B, how='right')
```

```
[27]:
```

	color	num	size
0	green	1.0	S
1	yellow	2.0	M
2	pink	NaN	L